

Implementing Oblivious Data Structures using Red-Black trees



UC SANTA CRUZ

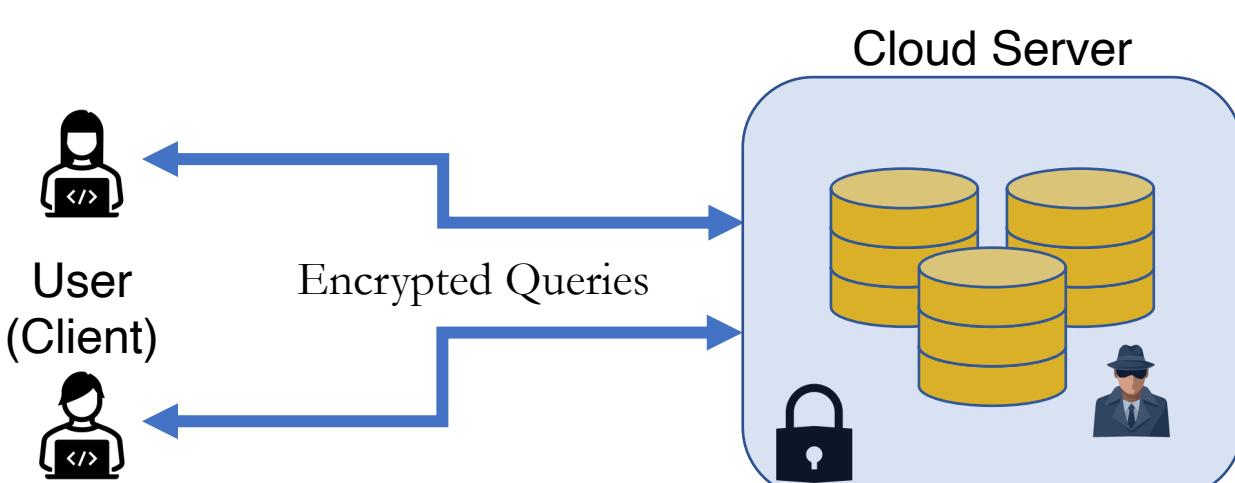
Prathamesh Pradeep Khole
University of California, Santa Cruz
pkhole@ucsc.edu
CSE-290X Course Project



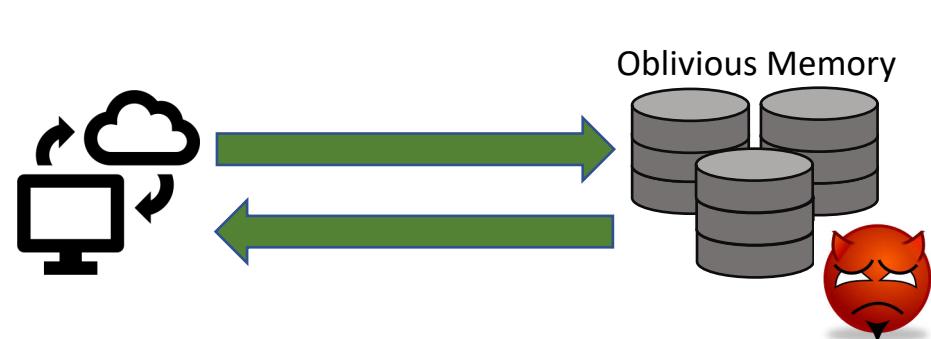
Baskin
Engineering

Introduction

- * Using cloud storage has become widespread nowadays.
- * The server or the cloud stores and protects the client's data.
- * The server is honest but also curious.



- * Since the data is encrypted, the only thing the server can learn about the data is the search and access pattern of the user.
- * Search and access patterns can reveal a lot about the data; as a result, data privacy may be compromised even in the presence of encryption.
- * Oblivious memory can help mitigate this problem.
- * Oblivious memory works by performing a certain number of dummy accesses with the actual ones to trick or confuse the observer.
- * Extra operations act like padding so that all operations appear to have accessed the same number of elements.
- * As a result, any two user operations are indistinguishable from one another for the observer.

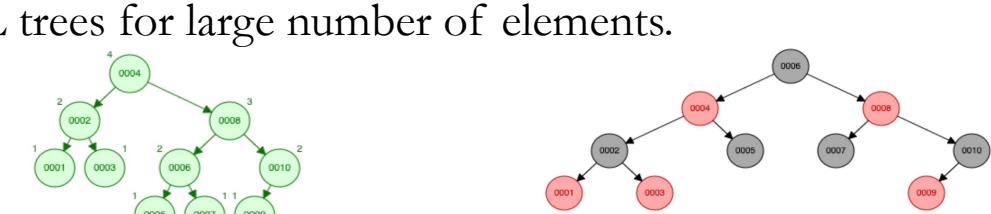


Problem Definition

- * This Project aims to implement efficient oblivious data structures using Red-Black trees with existing searchable encryption schemes.
- * To do this, we implement an Oblivious Random Access Machine (ORAM) using a Red-Black tree as the underlying data structure.
- * Current ORAM implementations also use balanced binary search trees, the most popular being the AVL trees.
- * AVL Trees are widely used in ORAMs since they are strictly balanced and thus provide fast lookups.
- * The reason behind using Red-Black trees for oblivious RAM is that Red-Black trees as data structures are faster at insertions and deletions than AVL trees.
- * Also, although Red-Black trees are less strictly balanced binary search trees, they have logarithmic search speeds.

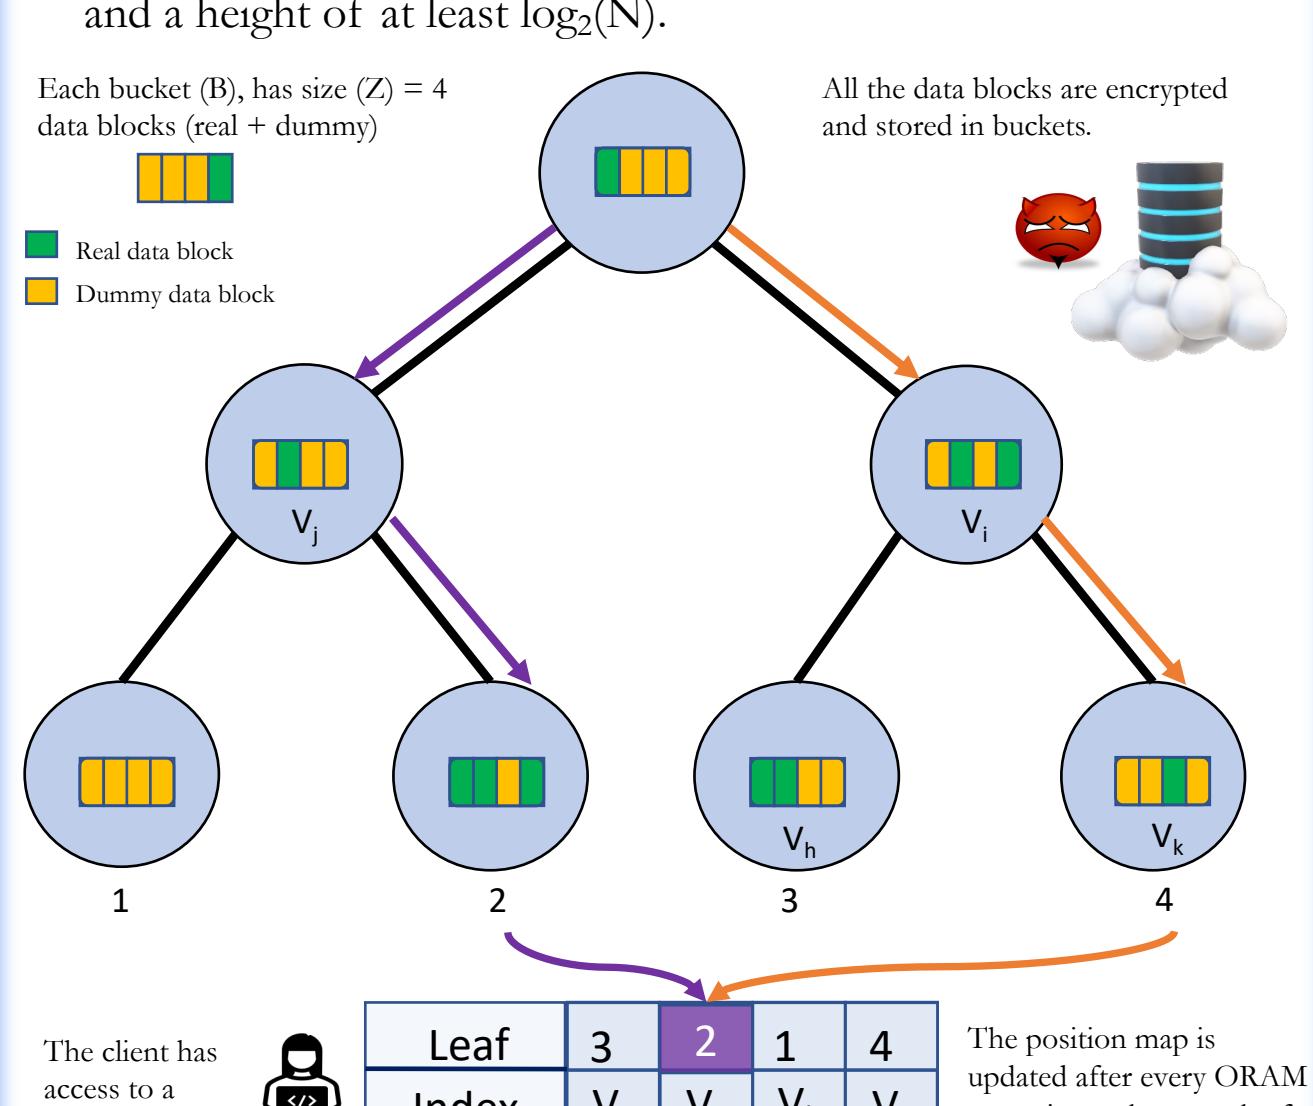
Observations and results:

- * Implemented ORAM with Red-Black tree as the underlying data structure which supports insertions, search and delete operations on encrypted data.
- * Experimentally verified the number of nodes accessed by Red-Black trees during insertion and delete operations is less than AVL trees for large number of elements.



Tree Based Oblivious RAM

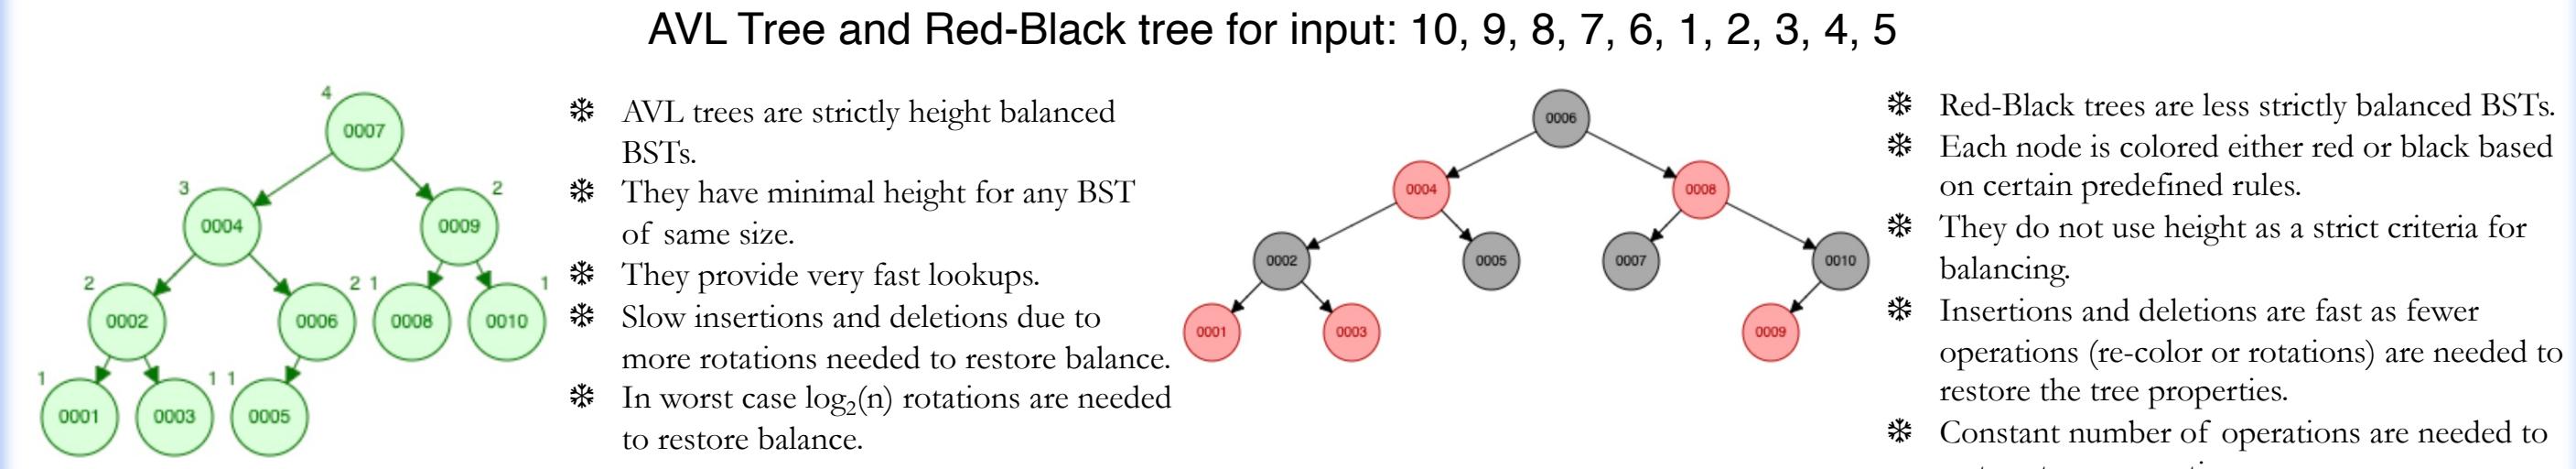
- * In case of tree based ORAMs the server stores the ORAM as a complete binary tree.
- * Complete binary tree is a type of binary tree where every level, excluding possibly last one is filled completely.
- * In case of ORAM, this binary tree has N leaves in the last level and a height of at least $\log_2(N)$.



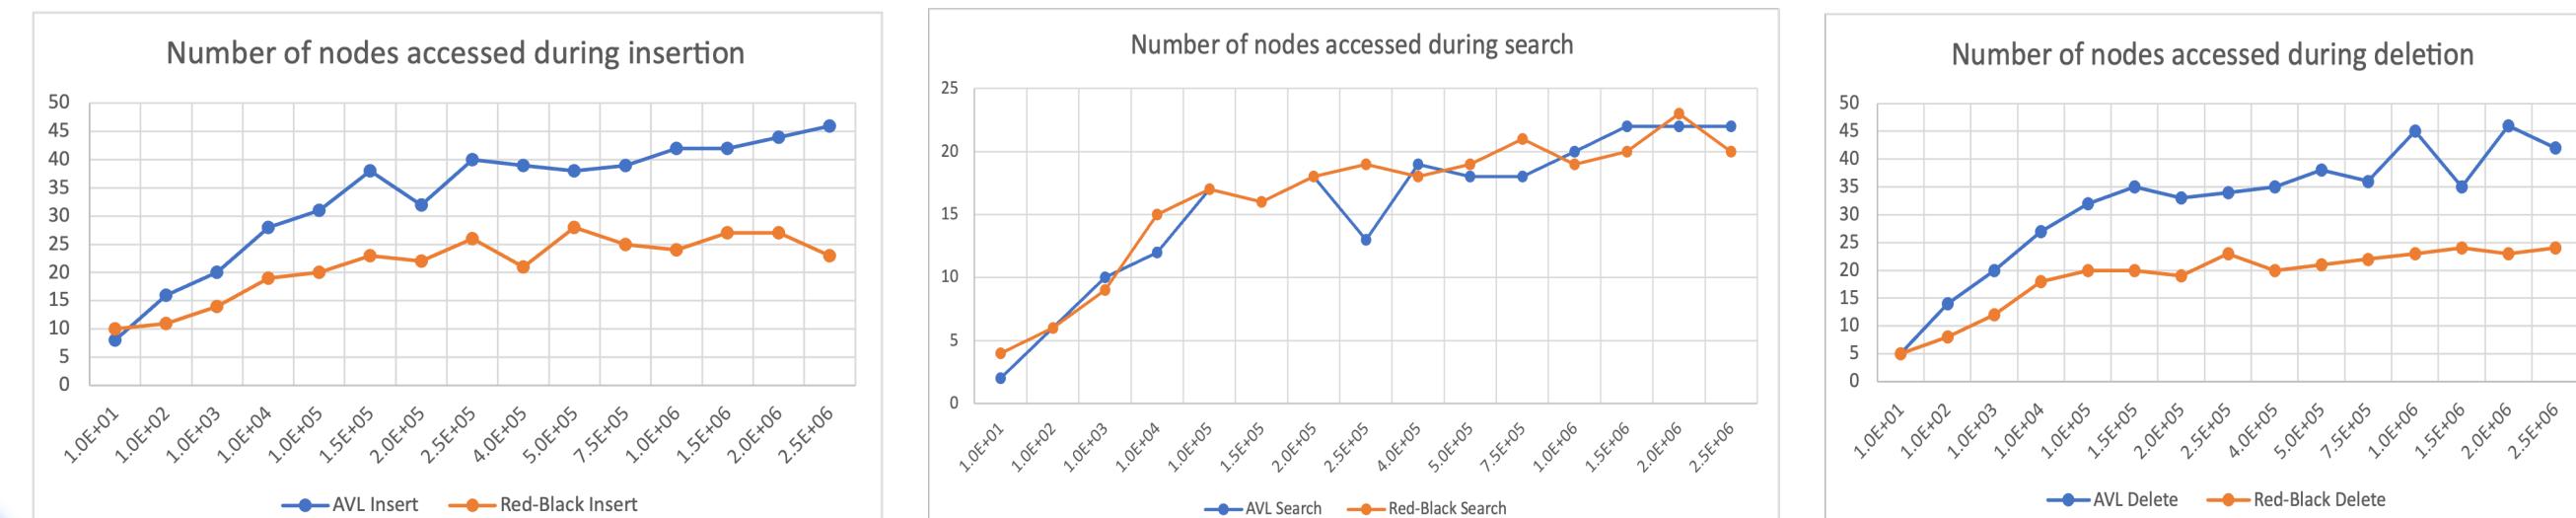
- * After each ORAM operation (e.g., read), the system logically maps the data block to a new random root->leaf path.
- * The eviction operation slowly pushes the re-mapped block from the root towards the selected leaf.

AVL Tree vs Red-Black Tree

- * AVL and Red-Black trees are self-balancing binary search trees but differ in maintaining the height balance.
- * Due to this they get appear different for same inputs.

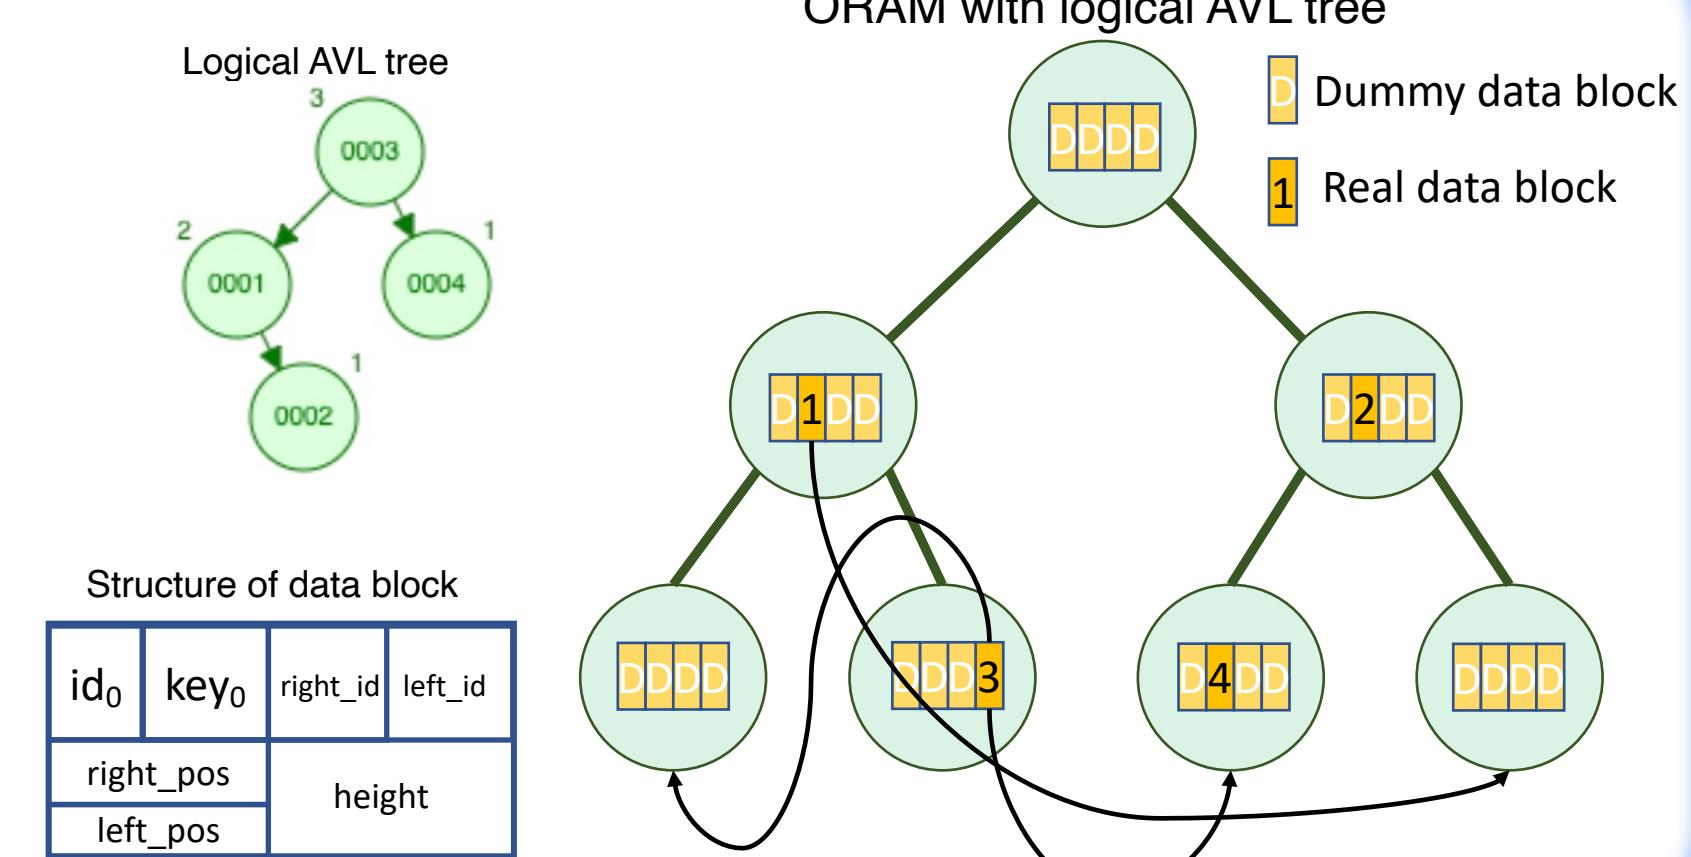


Comparison between AVL tree and Red-Black tree for insert, search and delete operations based on number of nodes accessed (y-axis) for n-elements (x-axis).



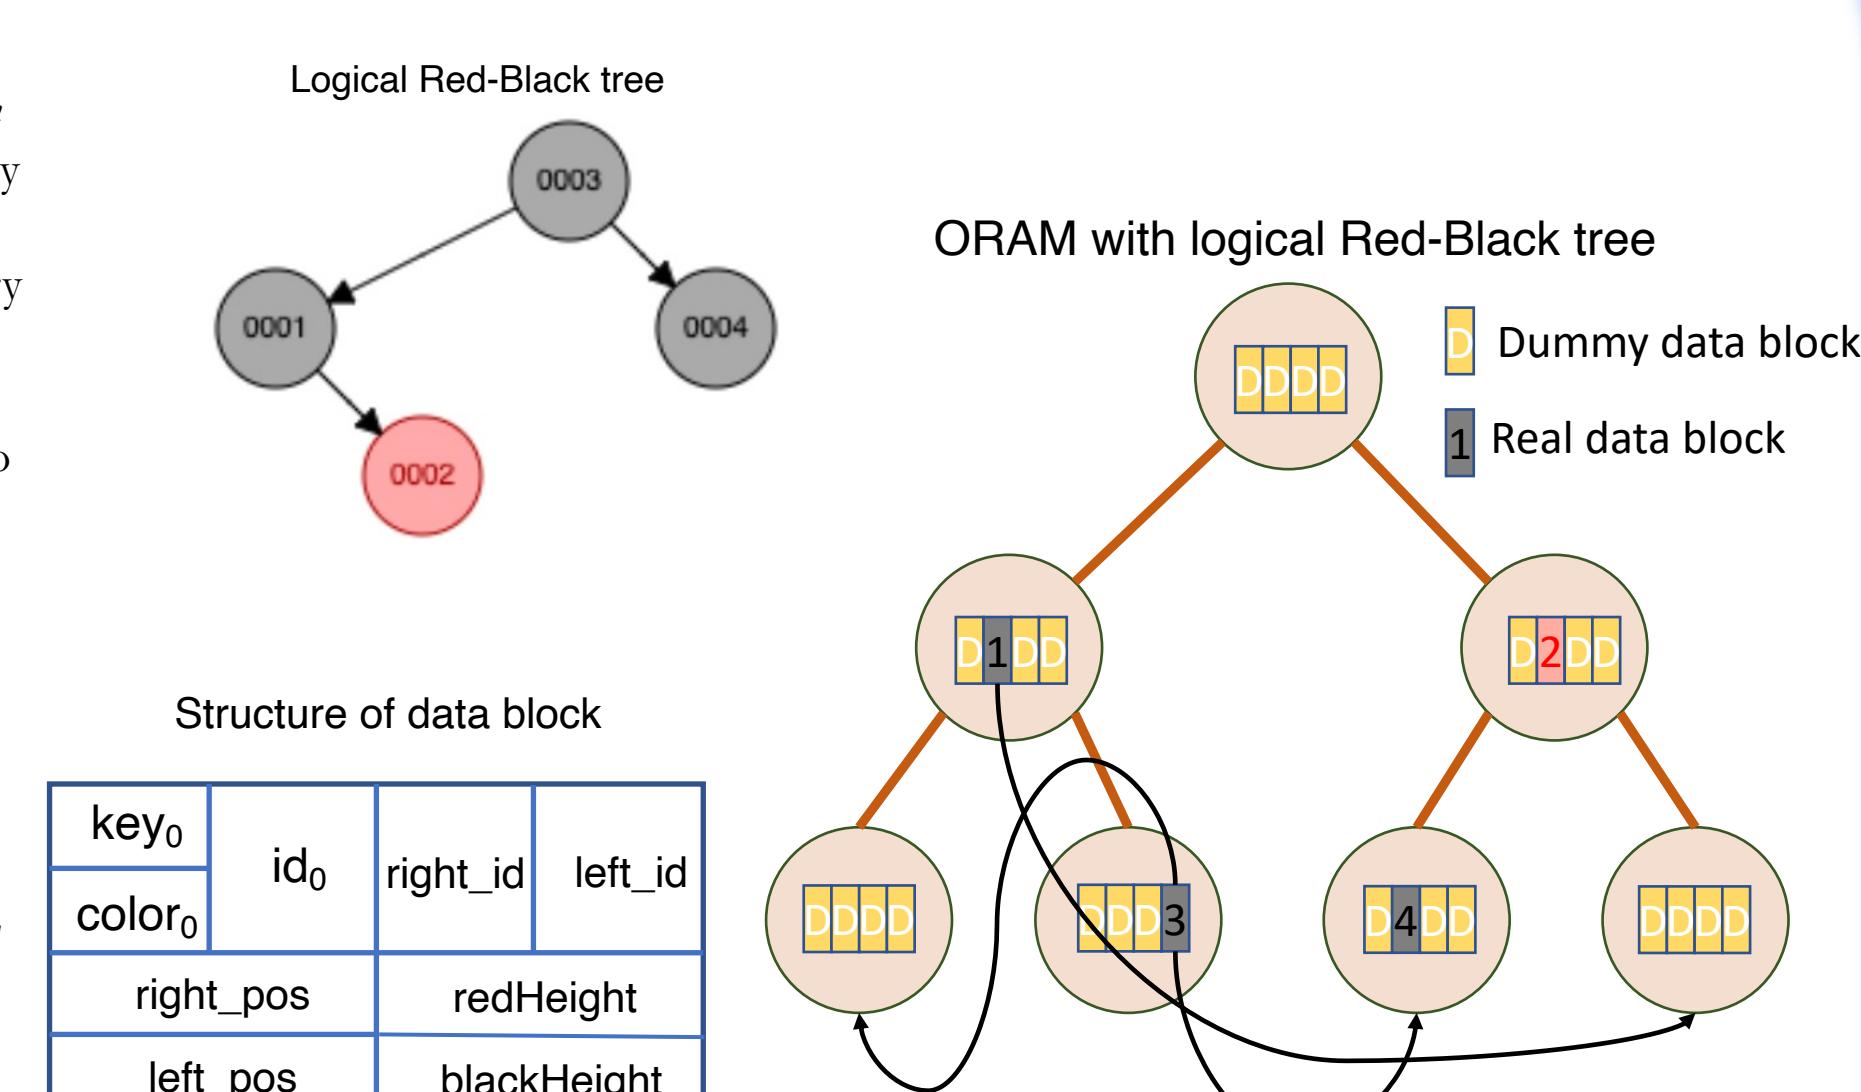
Related Work

- * The paper, Path ORAM first introduced Path ORAM as a practical and simple Oblivious RAM with a small amount of client storage.
- * Then Paper, *Oblivious Data Structures* discusses using logical Binary search trees on top of ORAM trees to speed up search process on the ORAM data.
- * AVL Trees are the most prevalent underlying logical data structures with ORAMs because of their fast lookups.
- * The paper, *Dynamic Searchable Encryption with Optimal Search in the Presence of Deletions* uses a Path ORAM with AVL trees as the underlying data structure.
- * This scheme makes use of a logical AVL tree built on top of an ORAM represented as a full binary tree having N leaf nodes.



Project's Approach

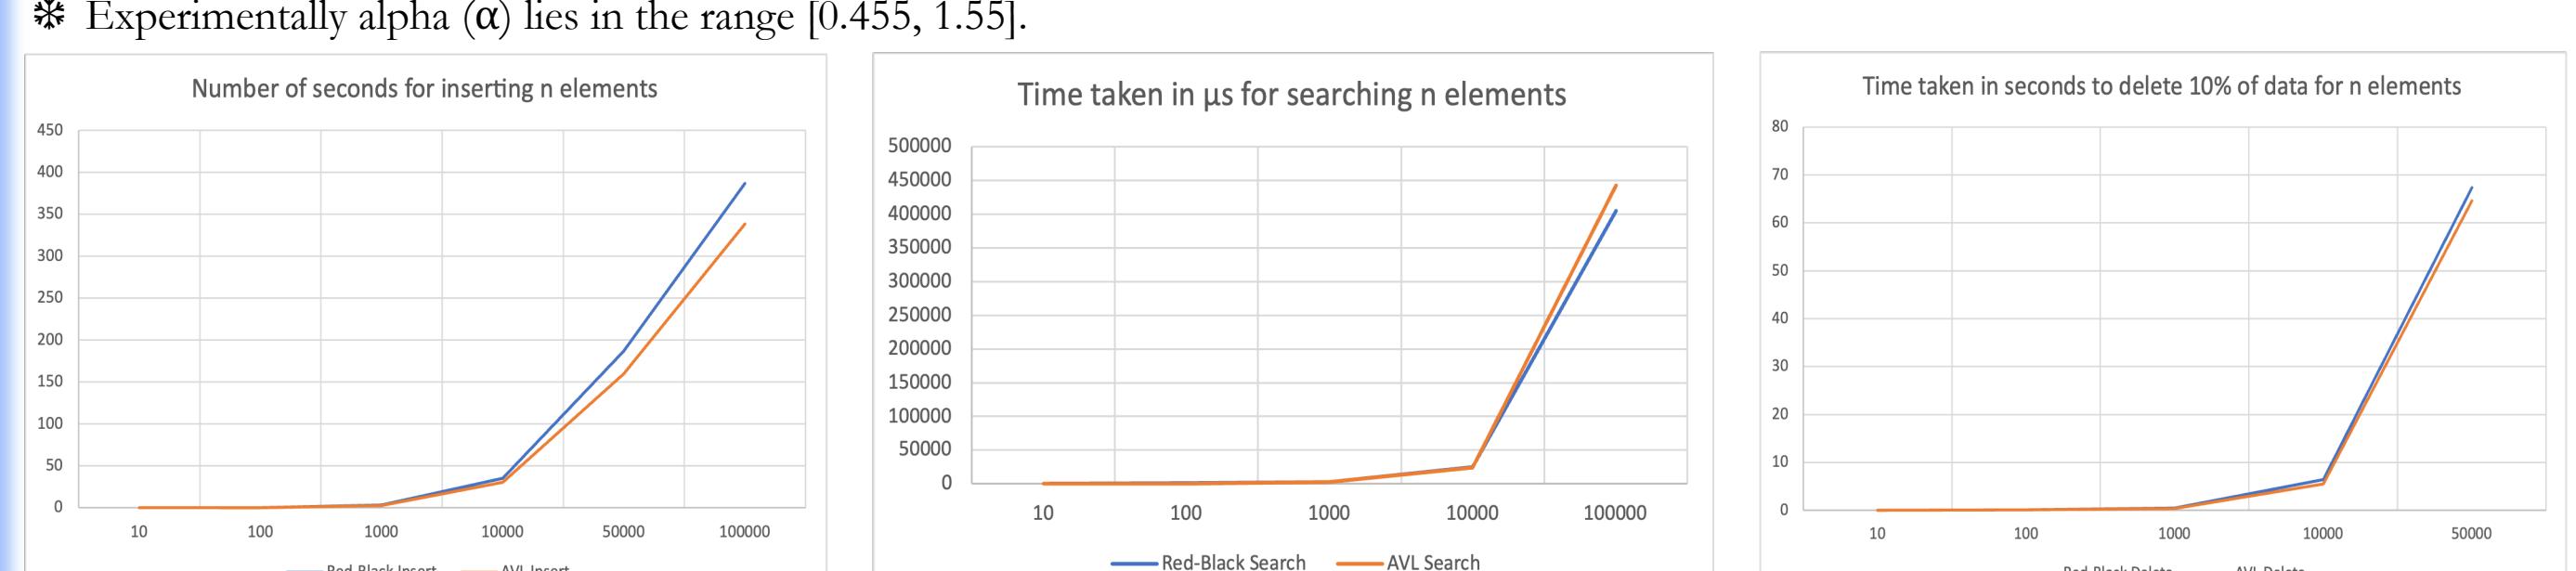
- * Our project aimed to improve the performance of *Dynamic Searchable Encryption with Optimal Search in the Presence of Deletions* by replacing the AVL tree data structure with a red-black tree, which is also a balanced binary search tree.
- * We leveraged the ORAM structure and updated the node and data block structure to support Red-Black trees, modifying the rotation and balance functions to handle the necessary re-coloring.
- * It is possible to implement these changes while still using the encryption schemes of the previous ORAM.
- * These changes allowed us to compare the performance of Red-Black trees versus AVL trees and demonstrate their potential to enhance the efficiency of this encryption technique.



Although the ORAM tree is like the previous implementations, the logical tree is different and thus uses different rules to maintain its properties.

Experimental Results

- * The AVL tree based ORAM performs slightly better than the new Red-Black tree-based approach.
- * We defined alpha (α) as an empirical parameter to find the worst-case number of nodes accessed during any operation in Red-Black trees based on the equation: worst-case nodes accessed = $2 * \alpha * \log_2(n)$, since worst case height of a Red-Black tree is $2 * \log_2(n+1)$ where n is the number of elements.
- * Experimentally alpha (α) lies in the range [0.455, 1.55].



Resources:

<https://www.cs.usfca.edu/~galles/visualization/RedBlack.html>, <https://www.cs.usfca.edu/~galles/visualization/AVLtree.html>, <https://eprint.iacr.org/2022/648.pdf>, <https://eprint.iacr.org/2013/280.pdf>, <https://eprint.iacr.org/2014/185.pdf>, <https://github.com/gharehchamani/OS-SSE>, <https://www.programiz.com/dsa/red-black-tree>, <https://www.programiz.com/dsa/avl-tree>, lecture slides of Professor Ioannis Demertzis, ODS slides by Amin Karbas.